

CS631 - Advanced Programming in the UNIX Environment

Department of Computer Science

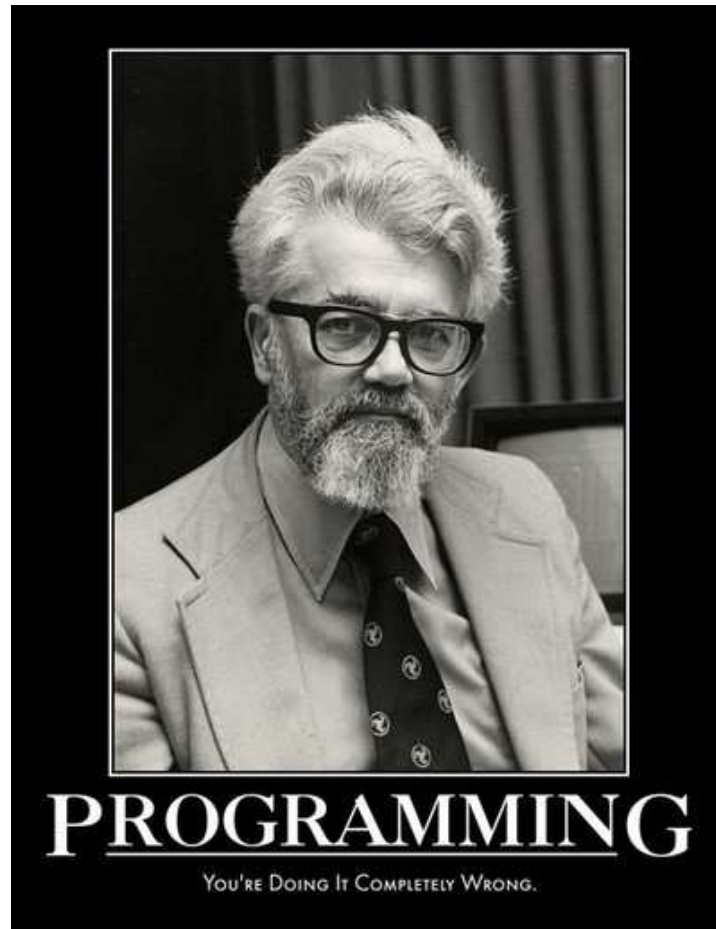
Stevens Institute of Technology

Jan Schaumann

`jschauma@stevens.edu`

`http://www.cs.stevens.edu/~jschauma/631/`

Programming



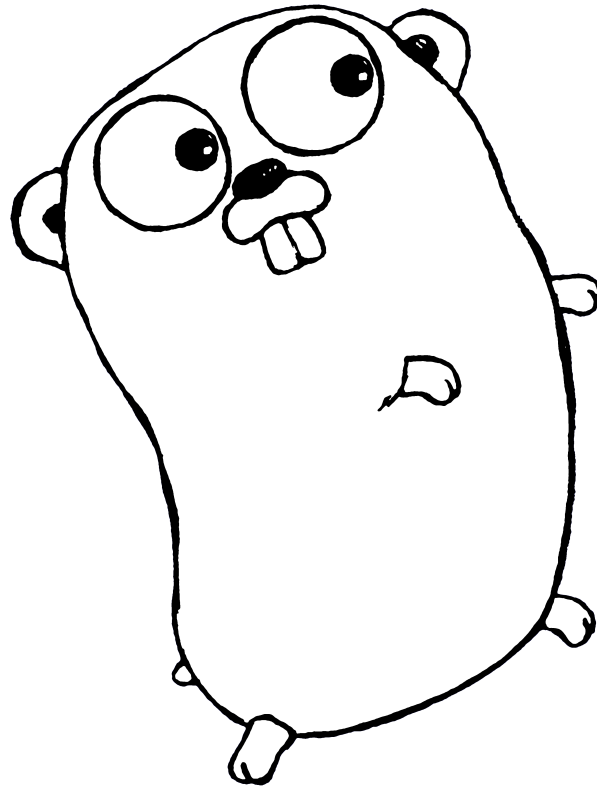
<http://is.gd/Z4CLky>

Programming



<http://cm.bell-labs.com/who/dmr/chist.html>

Extra Fun!



<http://golang.org>

In a nutshell: the "what"

```
$ ls /bin
```

```
[      csh      ed      ls      pwd      sleep
cat     date     expr    mkdir   rcmd     stty
chio    dd      hostname mt      rcp      sync
chmod   df      kill    mv      rm       systrace
cp      domainname ksh     pax     rmdir    tar
cpio    echo     ln      ps      sh       test
$
```

In a nutshell: the "what"

```
$ grep "(int" /usr/include/sys/socket.h
int accept(int, struct sockaddr * __restrict, socklen_t * __restrict);
int bind(int, const struct sockaddr *, socklen_t);
int connect(int, const struct sockaddr *, socklen_t);
int getsockopt(int, int, int, void * __restrict, socklen_t * __restrict);
int listen(int, int);
ssize_t recv(int, void *, size_t, int);
ssize_t recvfrom(int, void * __restrict, size_t, int,
ssize_t recvmsg(int, struct msghdr *, int);
ssize_t send(int, const void *, size_t, int);
ssize_t sendto(int, const void *,
ssize_t sendmsg(int, const struct msghdr *, int);
int setsockopt(int, int, int, const void *, socklen_t);
int socket(int, int, int);
int socketpair(int, int, int, int *);
$
```

In a nutshell: the "what"

- gain an understanding of the UNIX operating systems
- gain (systems) programming experience
- understand fundamental OS concepts (with focus on UNIX family):
 - multi-user concepts
 - basic and advanced I/O
 - process relationships
 - interprocess communication
 - basic network programming using a client/server model

In a nutshell: the "how"

```
static char dot[] = ".", *dotav[] = { dot, NULL };
struct winsize win;
int ch, fts_options;
int kflag = 0;
const char *p;

setprogname(argv[0]);
setlocale(LC_ALL, "");

/* Terminal defaults to -Cq, non-terminal defaults to -1. */
if (isatty(STDOUT_FILENO)) {
    if (ioctl(STDOUT_FILENO, TIOCGWINSZ, &win) == 0 &&
        win.ws_col > 0)
        termwidth = win.ws_col;
    f_column = f_nonprint = 1;
} else
    f_singlecol = 1;

/* Root is -A automatically. */
if (!getuid())
    f_listdot = 1;

fts_options = FTS_PHYSICAL;
while ((ch = getopt(argc, argv, "1ABCFLRSTWabcdfghiklmnopqrstuwX")) != -1) {
    switch (ch) {
        /*
         * The -1, -C, -l, -m and -x options all override each other so
         * shell aliasing works correctly.
         */
        case '1':
            f_singlecol = 1;
```


In a nutshell: the "how"

```
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
$ ./cmd
$ echo "Hooray!"
Hooray!
$
```

In a nutshell: the "how"

```
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
cmd.c: In function 'main':
cmd.c:19: error: parse error before "return"
$
```

In a nutshell: the "how"

```
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
cmd.c: In function 'main':
cmd.c:19: error: parse error before "return"
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
$ ./cmd
Memory fault (core dumped)
$
```

In a nutshell: the "how"

```
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
cmd.c: In function 'main':
cmd.c:19: error: parse error before "return"
$ $EDITOR cmd.c
$ cc -Wall -g -o cmd cmd.c
$ ./cmd
Memory fault (core dumped)
$ echo "!@#!@!!!??#@!"
!@#!@!!!??#@!
$ gdb ./cmd cmd.core
Program terminated with signal 11, Segmentation fault.
Loaded symbols for /usr/libexec/ld.elf_so
#0  0xbbbc676a in __findenv () from /usr/lib/libc.so.12
(gdb)
```

In a nutshell

The "why":

- understanding how UNIX works gives you insights in other OS concepts
- system level programming experience is invaluable as it forms the basis for most other programming and even *use* of the system
- system level programming in C helps you understand general programming concepts
- most higher level programming languages (eventually) call (or implement themselves) standard C library functions

About this class

Textbook:

- “Advanced Programming in the UNIX Environment”, by W. Richard Stevens, Stephen A. Rago (2nd Edition)

Help:

- <http://lists.stevens.edu/cgi-bin/mailman/listinfo/cs631apue>
- <https://twitter.com/#!/cs631apue>

Grading:

- 5 homework assignments, worth 20 points each
- 1 midterm project, worth 100 points
- 1 final exam, worth 100 points
- 1 final project, worth 200 points
- no curve

Syllabus

- 2013-08-26: Introduction, UNIX history, UNIX Programming Basics
- 2013-09-09: File I/O, File Sharing
- 2013-09-16: Files and Directories, Filesystems
- 2013-09-23: System Data Files, Time & Date, Process Environment
- 2013-09-30: Process Control, Signals
- 2013-10-07: Signals
- 2013-10-15: Interprocess Communication
- 2013-10-21: Advanced I/O: Nonblocking I/O, Polling, and Record Locking
- 2013-10-28: Daemon Processes, final project discussion
- 2013-11-04: UNIX tools: make(1), gdb(1), revision control, etc.
- 2013-11-11: Encryption
- 2013-11-18 – 2013-11-25: Code reading and discussions
- 2013-12-02: Final Exam

`linux-lab.cs.stevens.edu`

`http://tarantula.phy.stevens.edu/wiki/index.php/Linux_Lab`

UNIX History

UNIX history

http://www.unix.org/what_is_unix/history_timeline.html

<http://www.levenez.com/unix/>

- Originally developed in 1969 at Bell Labs by Ken Thompson and Dennis Ritchie.
- 1973, Rewritten in C. This made it portable and changed the history of OS
- 1974: Thompson, Joy, Haley and students at Berkeley develop the Berkeley Software Distribution (BSD) of UNIX
- two main directions emerge: BSD and what was to become “System V”

UNIX history

- 1984 4.2BSD released (TCP/IP)
- 1986 4.3BSD released (NFS)
- 1991 Linus Torvalds starts working on the Linux kernel
- 1993 Following the settlement of USL vs. BSDi: NetBSD, then FreeBSD are created
- 1994 Single UNIX Specification introduced
- 1995 4.4BSD-Lite Release 2 (last CSRG release); OpenBSD forked off NetBSD
- 2000 Darwin created (derived from NeXT, FreeBSD, NetBSD)
- ...

Some of today's main Unix versions:

- mostly BSD: *BSD, Linux, Mac OS X
- mostly SysV: Solaris, HP-UX, IRIX

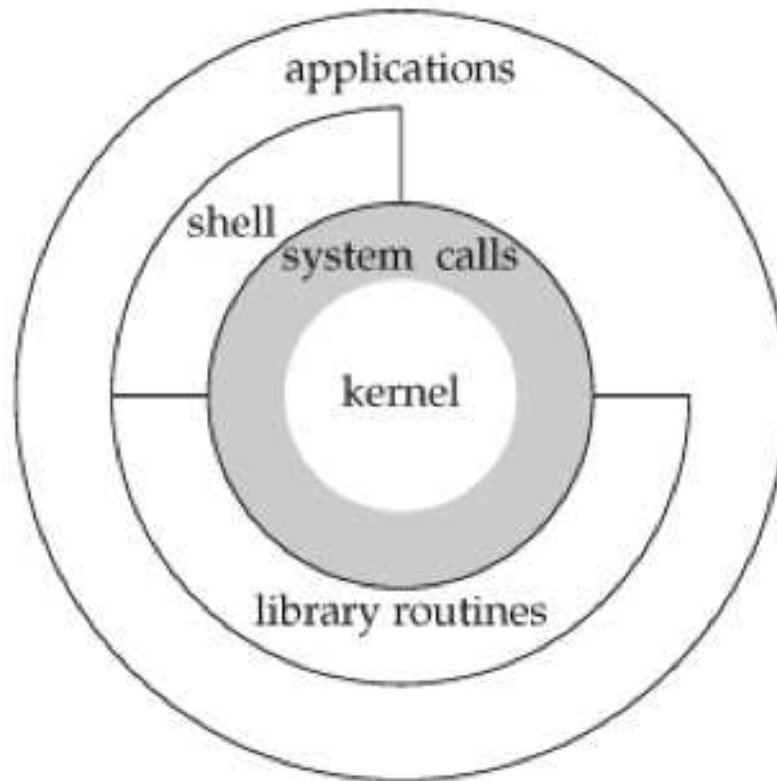
Some UNIX versions

More UNIX (some generic, some trademark, some just unix-like):

1BSD	2BSD	3BSD	4BSD	4.4BSD Lite 1
4.4BSD Lite 2	386 BSD	A/UX	Acorn RISC iX	AIX
AIX PS/2	AIX/370	AIX/6000	AIX/ESA	AIX/RT
AMiX	AOS Lite	AOS Reno	ArchBSD	ASV
Atari Unix	BOS	BRL Unix	BSD Net/1	BSD Net/2
BSD/386	BSD/OS	CB Unix	Chorus	Chorus/MiX
Coherent	CTIX	Darwin	Debian GNU/Hurd	DEC OSF/1 ACP
Digital Unix	DragonFly BSD	Dynix	Dynix/ptx	ekkoBSD
FreeBSD	GNU	GNU-Darwin	HPBSD	HP-UX
HP-UX BLS	IBM AOS	IBM IX/370	Interactive 386/ix	Interactive IS
IRIX	Linux	Lites	LSX	Mac OS X
Mac OS X Server	Mach	MERT	MicroBSD	Mini Unix
Minix	Minix-VMD	MIPS OS	MirBSD	Mk Linux
Monterey	more/BSD	mt Xinu	MVS/ESA OpenEdition	NetBSD
NeXTSTEP	NonStop-UX	Open Desktop	Open UNIX	OpenBSD
OpenServer	OPENSTEP	OS/390 OpenEdition	OS/390 Unix	OSF/1
PC/IX	Plan 9	PWB	PWB/UNIX	QNX
QNX RTOS	QNX/Neutrino	QUNIX	ReliantUnix	Rhapsody
RISC iX	RT	SCO UNIX	SCO UnixWare	SCO Xenix
SCO Xenix System V/386	Security-Enhanced Linux	Sinix	Sinix ReliantUnix	Solaris
SPIX	SunOS	Tru64 Unix	Trusted IRIX/B	Trusted Solaris
Trusted Xenix	TS	UCLA Locus	UCLA Secure Unix	Ultrix
Ultrix 32M	Ultrix-11	Unicos	Unicos/mk	Unicox-max
UNICS	UNIX 32V	UNIX Interactive	UNIX System III	UNIX System IV
UNIX System V	UNIX System V Release 2	UNIX System V Release 3	UNIX System V Release 4	UNIX System V/286
UNIX System V/386	UNIX Time-Sharing System	UnixWare	UNSW	USG
Venix	Wollogong	Xenix OS	Xinu	xMach

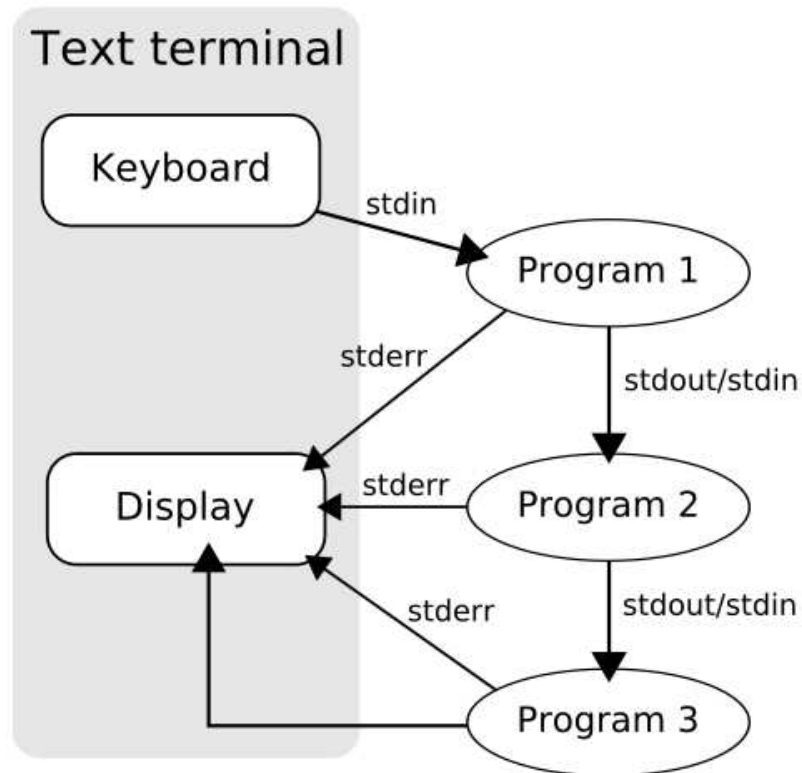
UNIX Basics

UNIX Basics: Architecture



UNIX Basics: Pipelines

Say "Thank you, Douglas McIlroy!"



<http://is.gd/vGH09J>

Program Design

“Consistency underlies all principles of quality.”
Frederick P. Brooks, Jr

Program Design

https://secure.wikimedia.org/wikipedia/en/wiki/Unix_philosophy

UNIX programs...

- ...are simple
- ...have a manual page
- ...follow the element of least surprise
- ...accept input from `stdin`
- ...generate output to `stdout`
- ...generate meaningful error messages to `stderr`
- ...have meaningful exit codes

Boot/Login process

```
[...]  
total memory = 768 MB  
avail memory = 732 MB  
timecounter: Timecounters tick every 10.000 msec  
mainbus0 (root)  
[...]  
boot device: xbd3  
root on xbd3a dumps on xbd3b  
mountroot: trying lfs...  
mountroot: trying ffs...  
root file system type: ffs  
init: copying out path '/sbin/init' 11  
[...]  
Starting local daemons:.  
Starting sendmail.  
Starting sshd.  
Starting snmpd.  
Starting cron.  
  
NetBSD/amd64 (panix.netmeister.org) (console)  
  
login:
```

Boot/Login process

```
[...]  
total memory = 768 MB  
avail memory = 732 MB  
timecounter: Timecounters tick every 10.000 msec  
mainbus0 (root)  
[...]  
boot device: xbd3  
root on xbd3a dumps on xbd3b  
mountroot: trying lfs...  
mountroot: trying ffs...  
root file system type: ffs  
init: copying out path '/sbin/init' 11  
[...]  
Starting local daemons:.  
Starting sendmail.  
Starting sshd.  
Starting snmpd.  
Starting cron.  
  
NetBSD/amd64 (panix.netmeister.org) (console)  
  
login: jschauma  
Password:
```

Boot/Login process

```
[...]  
total memory = 768 MB  
avail memory = 732 MB  
timecounter: Timecounters tick every 10.000 msec  
mainbus0 (root)
```

```
[...]  
boot device: xbd3  
root on xbd3a dumps on xbd3b  
mountroot: trying lfs...  
mountroot: trying ffs...  
root file system type: ffs  
init: copying out path '/sbin/init' 11  
[...]  
Starting local daemons:.  
Starting sendmail.  
Starting sshd.  
Starting snmpd.  
Starting cron.
```

```
NetBSD/amd64 (panix.netmeister.org) (console)
```

```
login: jschauma  
Password:  
Last login: Sat Sep 10 14:27:56 2011 on console  
Copyright (c) 1982, 1986, 1989, 1991, 1993  
The Regents of the University of California. All rights reserved.
```

```
NetBSD 5.0.2 (PANIX-VC) #2: Tue Oct 19 16:30:57 EDT 2010
```

```
Welcome to NetBSD!
```

```
$
```

Soooo... what exactly is a "shell"?

```
$ wget http://www.cs.stevens.edu/~jschauma/631/simple-shell.c
$ more simple-shell.c
$ cc -Wall -o mysh simple-shell.c
$ ./mysh
$$ /bin/ls
[...]
$$ ^D
$
```

Files and Directories

- The UNIX filesystem is a tree structure, with all partitions mounted under the root (/). File names may consist of any character except / and NUL as pathnames are a sequence of zero or more filenames separated by /'s.

Files and Directories

- The UNIX filesystem is a tree structure, with all partitions mounted under the root (/). File names may consist of any character except / and NUL as pathnames are a sequence of zero or more filenames separated by /'s.
- Directories are special "files" that contain mappings between *inodes* and *filenames*, called directory entries.

Files and Directories

- The UNIX filesystem is a tree structure, with all partitions mounted under the root (/). File names may consist of any character except / and NUL as pathnames are a sequence of zero or more filenames separated by /'s.
- Directories are special "files" that contain mappings between *inodes* and *filenames*, called directory entries.
- All processes have a current working directory from which all relative paths are specified. (Absolute paths begin with a slash, relative paths do not.)

Listing files in a directory

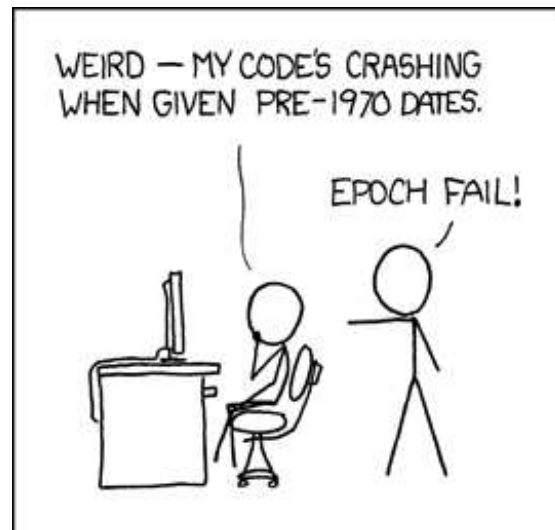
```
$ wget http://www.cs.stevens.edu/~jschauma/631/simple-ls.c
$ more simple-ls.c
$ cc -Wall -o myls simple-ls.c
$ ./mysls .
[...]
$
```

User Identification

- *User IDs* and *group IDs* are numeric values used to identify users on the system and grant permissions appropriate to them.
- *Group IDs* come in two types; *primary* and *secondary*.

Unix Time Values

Calendar time: measured in seconds since the UNIX epoch (Jan 1, 00:00:00, 1970, GMT). Stored in a variable of type `time_t`.



<https://www.xkcd.com/376/>

https://secure.wikimedia.org/wikipedia/en/wiki/Year_2038_problem

Unix Time Values

Process time: central processor resources used by a process. Measured in *clock ticks* (`clock_t`). Three values:

- clock time
- user CPU time
- system CPU time

```
$ time grep -r _POSIX_SOURCE /usr/include >/dev/null
```

System Calls and Library Functions, Standards

System Calls and Library Functions

- *System calls* are entry points into kernel code where their functions are implemented. Documented in section 2 of the manual (e.g. `write(2)`).
- *Library calls* are transfers to user code which performs the desired functions. Documented in section 3 of the manual (e.g. `printf(3)`).

Standards

- ANSI C (X3.159-1989) C89, C9X/C99 (ISO/IEC 9899), C11 (ISO/IEC 9899:2011)
- IEEE POSIX (1003.1-2008) / SUSv4

Standard I/O

- Standard I/O:
 - file descriptors: Small, non-negative integers which identify a file to the kernel. The shell can redirect any file descriptor.
 - kernel provides **unbuffered** I/O through e.g. `open read write lseek close`
 - kernel provides **buffered** I/O through e.g. `getc putc fopen fread fwrite`

Standard I/O: cat(1)

```
$ for f in cat cat2; do
    wget http://www.cs.stevens.edu/~jschauma/631/simple-${f}.c
    cc -Wall -o my${f} simple-${f}.c
    ./my${f} <simple-${f}.c
done
[...]
$
```

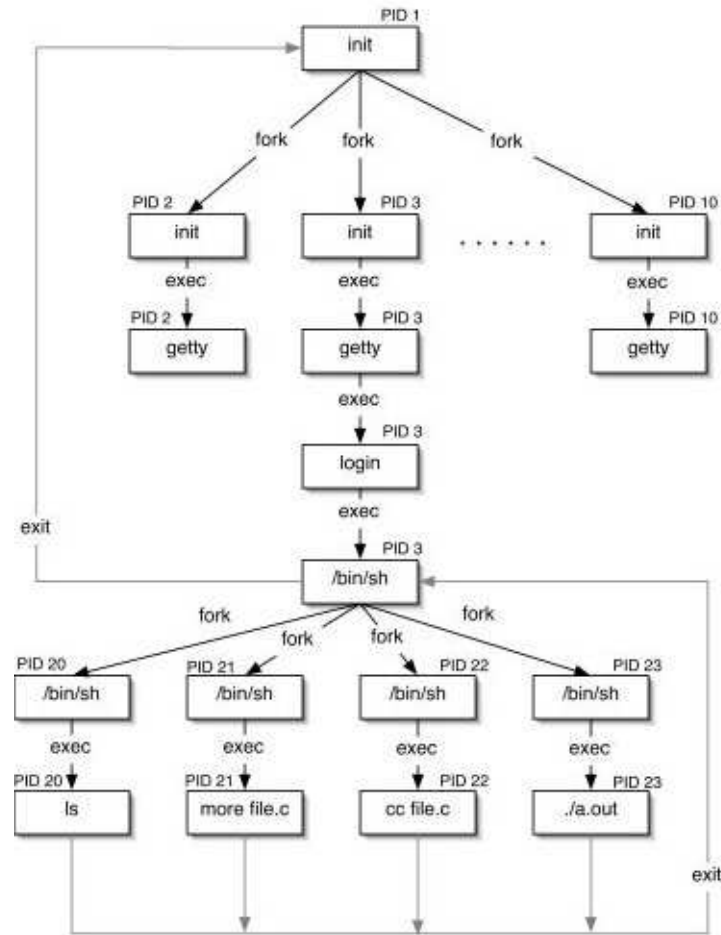
Processes

Programs executing in memory are called *processes*.

- Programs are brought into memory via one of the six `exec(3)` functions. Each process is identified by a guaranteed unique non-negative integer called the *processes ID*. New processes can only be created via the `fork(2)` system call.
- process control is performed mainly by the `fork(2)`, `exec(3)` and `waitpid(2)` functions.

```
$ wget http://www.cs.stevens.edu/~jschauma/631/pid.c
$ more pid.c
$ cc -Wall -o mypid pid.c
$ ./mypid .
[...]
$ echo $$
[...]
```


Processes



Processes

```
$ pstree -hapun | more
```

Signals

- Signals notify a *process* that a condition has occurred. Signals may be
 - ignored
 - allowed to cause the default action
 - caught and control transferred to a user defined function

```
$ wget http://www.cs.stevens.edu/~jschauma/631/simple-shell2.c
$ more simple-shell2.c
$ cc -Wall -o mysh simple-shell2.c
$ ./mysh
$$ /bin/ls
[...]
$$ ^C
Caught SIGINT!
```

Important ANSI C Features, Error Handling

- Important ANSI C Features:

- function prototypes
- generic pointers (`void *`)
- abstract data types (e.g. `pid_t`, `size_t`)

- Error Handling:

- meaningful return values
- `errno` variable
- look up constant error values via two functions:

```
#include <string.h>
char *strerror(int errnum)
```

Returns: pointer to message string

```
#include <stdio.h>
void perror(const char *msg)
```

Homework

- read `intro(2)`, Stevens 1 & 2
- follow, test and understand all examples from this lecture
- sign up on the course mailing list (using a Stevens email address)
- bookmark these websites:
 - <http://www.cs.stevens.edu/~jschauma/631/>
 - <http://pubs.opengroup.org/onlinepubs/9699919799/>
- ensure you have an account on `linux-lab.cs.stevens.edu`; see http://tarantula.phy.stevens.edu/wiki/index.php/Linux_Lab
- <http://www.cs.stevens.edu/~jschauma/631/f13-intro.html>
- <http://www.cs.stevens.edu/~jschauma/631/f13-hw1.html>