

# CS631 - Advanced Programming in the UNIX Environment

—

## HTTP, Dæmon processes, System Logging, Shared Libraries

---

Department of Computer Science  
Stevens Institute of Technology  
Jan Schaumann

`jschauma@stevens.edu`

`http://www.cs.stevens.edu/~jschauma/631/`

# HTTP

---

`http://www.cs.stevens.edu/~jschauma/631/f13-final-project.html`

## Hypertext Transfer Protocol

RFC2616

# HTTP

---

HTTP is a request/response protocol.

# The Hypertext Transfer Protocol

---

HTTP is a request/response protocol:

1. client sends a request to the server
2. server responds

# The Hypertext Transfer Protocol

---

HTTP is a request/response protocol:

1. client sends a request to the server

- request method
- URI
- protocol version
- request modifiers
- client information

2. server responds

## HTTP: A client request

---

```
$ telnet www.google.com 80
Trying 2607:f8b0:400c:c02::93...
Connected to www.google.com.
Escape character is '^]'.
GET / HTTP/1.0
```

# The Hypertext Transfer Protocol

---

HTTP is a request/response protocol:

1. client sends a request to the server
  - request method
  - URI
  - protocol version
  - request modifiers
  - client information
2. server responds
  - status line (including success or error code)
  - server information
  - entity metainformation
  - content

## HTTP: a server response

---

HTTP/1.0 200 OK

Date: Mon, 22 Oct 2012 03:08:18 GMT

Content-Type: text/html; charset=ISO-8859-1

Server: gws

```
<!doctype html><html itemscope="itemscope"
itemtype="http://schema.org/WebPage"><head><meta content="Search the
world's information, including webpages, images, videos and more. Google
has many special features to help you find exactly what you're looking
for." name="description"><meta content="noodp" name="robots"><meta
itemprop="image"
content="/images/google_favicon_128.png"><title>Google</title><script>
window.google={kEI:"oriEUNmMGMX50gH6kYGwBw",getEI:function(a){var
b;while(a&&!(a.getAttribute&&(b=a.getAttribute("eid"))))a=a.parentNode;return
b||google.kEI},https:function(){return
window.location.protocol=="https:"},kEXPI:"25657,30316,39523,39977,40362
```



# The Hypertext Transfer Protocol

---

Server status codes:

- 1xx – Informational; Request received, continuing process
- 2xx – Success; The action was successfully received, understood, and accepted
- 3xx – Redirection; Further action must be taken in order to complete the request
- 4xx – Client Error; The request contains bad syntax or cannot be fulfilled
- 5xx – Server Error; The server failed to fulfill an apparently valid request

## HTTP: A client request

---

```
$ telnet www.cs.stevens.edu 80
Trying 155.246.89.84...
Connected to tarantula.srcit.stevens-tech.edu.
Escape character is '^]'.
GET / HTTP/1.0
```

```
HTTP/1.1 200 OK
Date: Mon, 04 Apr 2011 02:16:14 GMT
Server: Apache/2.2.9 (Debian) DAV/2 SVN/1.5.1 PHP/5.2.6-1+lenny9 with Suhosin-Patch m
Last-Modified: Wed, 17 Nov 2010 19:25:54 GMT
Content-Type: text/html
```

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.0 Transitional//EN">
<html>
<head>
<title>SRCIT wiki page</title>
<meta http-equiv="REFRESH"
content="0;url=http://www.srcit.stevens.edu/wiki"></HEAD>
<BODY>
</BODY>
</HTML>
```

## HTTP - more than just text

---

HTTP is a *Transfer Protocol* – serving *data*, not any specific text format.

- Accept-Encoding client header can specify different formats such as *gzip*, *Shared Dictionary Compression over HTTP (SDCH)* etc.
- corresponding server headers: Content-Type and Content-Encoding



## HTTP - more than just static data

---

HTTP is a *Transfer Protocol* – what is transferred need not be static; resources may generate different data to return based on many variables.

- CGI – resource is *executed*, needs to generate appropriate response headers
- server-side scripting (ASP, PHP, Perl, ...)
- client-side scripting (JavaScript/ECMAScript/JScript,...)
- applications based on HTTP, using:
  - AJAX
  - RESTful services
  - JSON, XML, YAML to represent state and abstract information

## Writing a *simple* HTTP server

---

- parse command-line options, initialize world, ...
- open socket
- run as a daemon, loop forever
  - accept connection
  - fork child to handle request
- upon SIGHUP re-read configuration, restart

## Writing a *simple* HTTP server

---

Processing requests consists of:

- reading request from socket
- parsing request
  - valid syntax?
  - type of request (GET, HEAD, POST)?
  - determine pathname
    - ~ translation
    - translate relative into absolute pathname
- generate server status response
- handle request

## Writing a *simple* HTTP server

---

Processing requests consists of:

- handling regular file request
  - stat(2) file
  - open(2) file
  - read(2) file
  - write(2) to socket
  - close(2) file
  - terminate connection
  - exit child handler
- handling CGI execution
  - setup environment
  - setup filedescriptors (stdin/stdout)
  - fork-exec executable

## Client-Server Model

---

- two categories of servers
  1. iterative
  2. concurrent



## Client-Server Model

---

- two categories of servers
  - 1. iterative
    - 1.1. wait for client request to arrive
    - 1.2. process the client request
    - 1.3. send the response back to the client
    - 1.4. go back to 1.1
  - 2. concurrent

## Client-Server Model

---

- two categories of servers

- 1. iterative

- 1.1. wait for client request to arrive

- 1.2. process the client request

- 1.3. send the response back to the client

- 1.4. go back to 1.1

- 2. concurrent

- 2.1. wait for client request to arrive

- 2.2. start a new server to handle this client's request

- 2.3. go back to 2.1

# Dæmon processes

---

So... what's a dæmon process anyway?



## Dæmon characteristics

---

Commonly, dæmon processes are created to offer a specific service.

Dæmon processes usually

- live for a long time
- are started at boot time
- terminate only during shutdown
- have no controlling terminal



## Dæmon characteristics

---

The previously listed characteristics have certain implications:

- do one thing, and one thing only
- no (or only limited) user-interaction possible
- consider current working directory
- how to create (debugging) output



## Writing a dæmon

---

- fork off the parent process
- change file mode mask (umask)
- create a unique Session ID (SID)
- change the current working directory to a safe place
- close (or redirect) standard file descriptors
- open any logs for writing
- enter actual dæmon code



## Writing a dæmon

---

```
int
daemon(int nochdir, int noclose)
{
    int fd;

    switch (fork()) {
    case -1:
        return (-1);
    case 0:
        break;
    default:
        _exit(0);
    }

    if (setsid() == -1)
        return (-1);

    if (!nochdir)
        (void)chdir("/");

    if (!noclose && (fd = open(_PATH_DEVNULL, O_RDWR, 0)) != -1) {
        (void)dup2(fd, STDIN_FILENO);
        (void)dup2(fd, STDOUT_FILENO);
        (void)dup2(fd, STDERR_FILENO);
        if (fd > STDERR_FILENO)
            (void)close(fd);
    }
    return (0);
}
```

## Dæmon conventions

---

- prevent against multiple instances via a *lockfile*
- allow for easy determination of PID via a *pidfile*
- configuration file convention `/etc/name.conf`
- include a system initialization script (for `/etc/rc.d/` or `/etc/init.d/`)
- re-read configuration file upon SIGHUP





# Logging

---

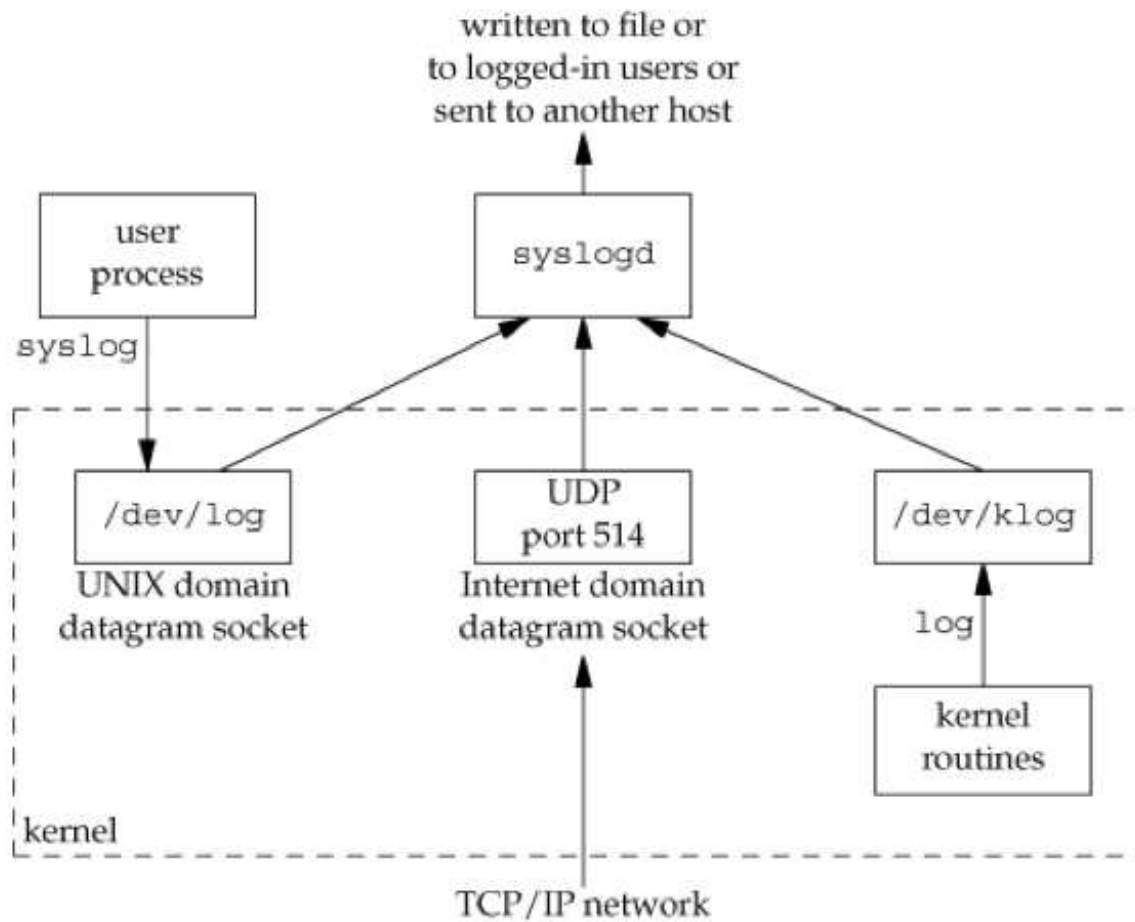
## A central logging facility

---

There are three ways to generate log messages:

- via the kernel routine `log(9)`
- via the userland routine `syslog(3)`
- via UDP messages to port 514

## A central logging facility



## syslog(3)

---

```
#include <syslog.h>

void openlog(const char *ident, int logopt, int facility);
void syslog(int priority, const char *message, ...);
```

openlog(3) allows us to set specific options when logging:

- prepend *ident* to each message
- specify logging options (LOG\_CONS | LOG\_NDELAY | LOG\_PERRO | LOG\_PID)
- specify a *facility* (such as LOG\_DAEMON, LOG\_MAIL etc.)

syslog(3) writes a message to the system message logger, tagged with *priority*.

A *priority* is a combination of a *facility* (as above) and a *level* (such as LOG\_DEBUG, LOG\_WARNING or LOG\_EMERG).

## Shared Libraries

---

What is a shared library, anyway?

## Shared Libraries

---

What is a shared library, anyway?

- contains a set of callable C functions (ie, implementation of function prototypes defined in `.h` header files)
- code is position-independent (ie, code can be executed anywhere in memory)
- shared libraries can be loaded/unloaded at execution time or at will
- libraries may be *static* or *dynamic*

## Shared Libraries

---

What is a shared library, anyway?

- contains a set of callable C functions (ie, implementation of function prototypes defined in `.h` header files)
- code is position-independent (ie, code can be executed anywhere in memory)
- shared libraries can be loaded/unloaded at execution time or at will
- libraries may be *static* or *dynamic*

```
$ man 3 fprintf
```

```
$ grep " fprintf" /usr/include/stdio.h
```

## Shared Libraries

---

How do shared libraries work?


- contents of *static* libraries are pulled into the executable at link time
- contents of *dynamic* libraries are used to resolve symbols at link time, but loaded at execution time by the *dynamic linker*
- contents of *dynamic* libraries may be loaded at any time via explicit calls to the dynamic linking loader interface functions



## Executable and Linkable Format

---

ELF is a file format for executables, object code, shared libraries etc.



More details: <http://www.cs.stevens.edu/~jschauma/631/elf.html>

<http://www.thegeekstuff.com/2012/07/elf-object-file-format/>

## Understanding object files

---

```
$ cc -Wall -c ldtest1.c ldtest2.c main.c
$ readelf -h ldtest1.o
[...]
$ cc *.o
$ readelf -h a.out
[...]
$ ldd a.out
[...]
$ readelf -h /lib/libc.so.6
[...]
$ readelf -s a.out | more
[...]
$ objdump -d -j .text a.out | more
[...]
$ nm -D a.out | more
[...]
$
```

## Statically Linked Shared Libraries

---

Static libraries:

- created by `ar(1)`
- usually end in `.a`
- contain a symbol table within the archive (see `ranlib(1)`)

## Statically Linked Shared Libraries

---

```
$ cc -Wall -c ldtest1.c
$ cc -Wall -c ldtest2.c
$ cc -Wall main.c
[...]
$ cc -Wall main.c ldtest1.o ldtest2.o
$
```

## Statically Linked Shared Libraries

---

```
$ cc -Wall -c ldtest1.c ldtest2.c
$ ar -vq libldtest.a ldtest1.o ldtest2.o
$ ar -t libldtest.a
$ cc -Wall main.c libldtest.a

$ cc -Wall -c main.c
$ cc main.o -L. -lldtest -o a.out.dyn
$ cc -static main.o -L. -lldtest -o a.out.static
$ ls -l a.out.*
$ ldd a.out.*
$ nm a.out.dyn | wc -l
$ nm a.out.static | wc -l
```

## Dynamically Linked Shared Libraries

---

Explicit loading of shared libraries:

- `dlopen(3)` creates a handle for the given library
- `dlsym(3)` returns the address of the given symbol
- 

```
$ cc -Wall setget.c
```

```
$ cc -Wall -rdynamic dlopenex.c -ldl
```

```
$ ./a.out
```

## Dynamically Linked Shared Libraries

---

Dynamic libraries:

- created by the compiler/linker (ie multiple steps)
- usually end in `.so`
- frequently have multiple levels of symlinks providing backwards compatibility / ABI definitions

## Dynamically Linked Shared Libraries

---

```
$ rm *.o libldtest*
$ cc -Wall -c -fPIC ldtest1.c
$ cc -Wall -c -fPIC ldtest2.c
$ mkdir lib
$ cc -shared -Wl,-soname,libldtest.so.1 -o lib/libldtest.so.1.0 ldtest1.o ldtest2.o
$ ln -s libldtest.so.1.0 lib/libldtest.so.1
$ ln -s libldtest.so.1.0 lib/libldtest.so
$ cc -static -Wall main.o -L./lib -lldtest
[...]
$ cc -Wall main.o -L./lib -lldtest
[...]
$ ./a.out
[...]
$ ldd a.out
[...]
```



## Dynamically Linked Shared Libraries

---

Wait, what?

```
$ export LD_LIBRARY_PATH=${LD_LIBRARY_PATH}:/lib
$ ldd a.out
[...]
$ ./a.out
[...]
$ mkdir lib2
$ cc -Wall -c -fPIC ldtest1.2.c
$ cc -shared -Wl,-soname,libldtest.so.1 -o lib2/libldtest.so.1.0 ldtest1.2.o ldtest2.o
$ ln -s libldtest.so.1.0 lib2/libldtest.so.1
$ ln -s libldtest.so.1.0 lib2/libldtest.so
$ export LD_LIBRARY_PATH=./lib2:$LD_LIBRARY_PATH
$ ldd a.out # note: no recompiling!
[...]
$ ./a.out
[...]
```

## Dynamically Linked Shared Libraries

---

Avoiding LD\_LIBRARY\_PATH:

```
$ cc -Wall main.o -L./lib -lldtest -Wl,-rpath,./lib
$ echo $LD_LIBRARY_PATH
[...]
$ ldd a.out
[...]
$ ./a.out
[...]
$ unset LD_LIBRARY_PATH
$ ldd a.out
[...]
$ ./a.out
[...]
$
```

## Dynamically Linked Shared Libraries

---

But:

```
$ cc -Wall -fPIC -c evil.c
$ cc -shared -Wl,-soname,libldtest.so.1 -o lib3/libldtest.so.1.0 \
    ldtest1.o ldtest2.o evil.o
$ export LD_PRELOAD=./lib3/libldtest.so.1.0
$ ldd a.out
[...]
$ ./a.out
[...]
$
```

## Dynamically Linked Shared Libraries

---

```
$ export LD_DEBUG=help # glibc>=2.1 only
$ ./a.out
[...]
$ LD_DEBUG=all ./a.out
[...]
```

## Homework

---

Same as last week:

<http://www.cs.stevens.edu/~jschauma/631/f13-hw3.html>

<http://www.cs.stevens.edu/~jschauma/631/f13-final-project.html>