

CS631 - Advanced Programming in the UNIX Environment

—

UNIX development tools

Department of Computer Science
Stevens Institute of Technology
Jan Schaumann

`jschauma@stevens.edu`

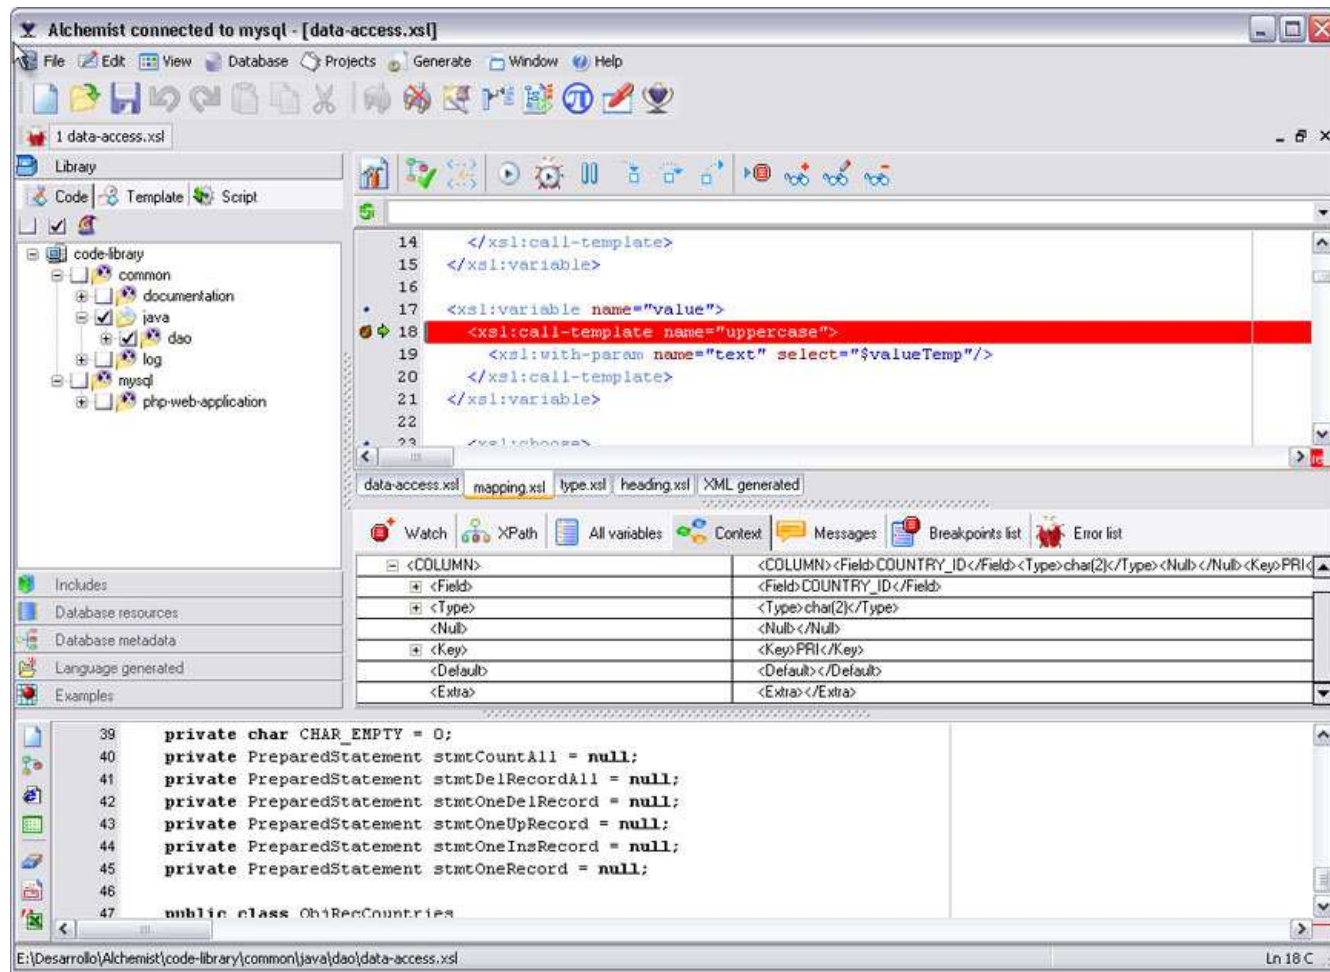
`http://www.cs.stevens.edu/~jschauma/631/`

Amended Final Project Requirements

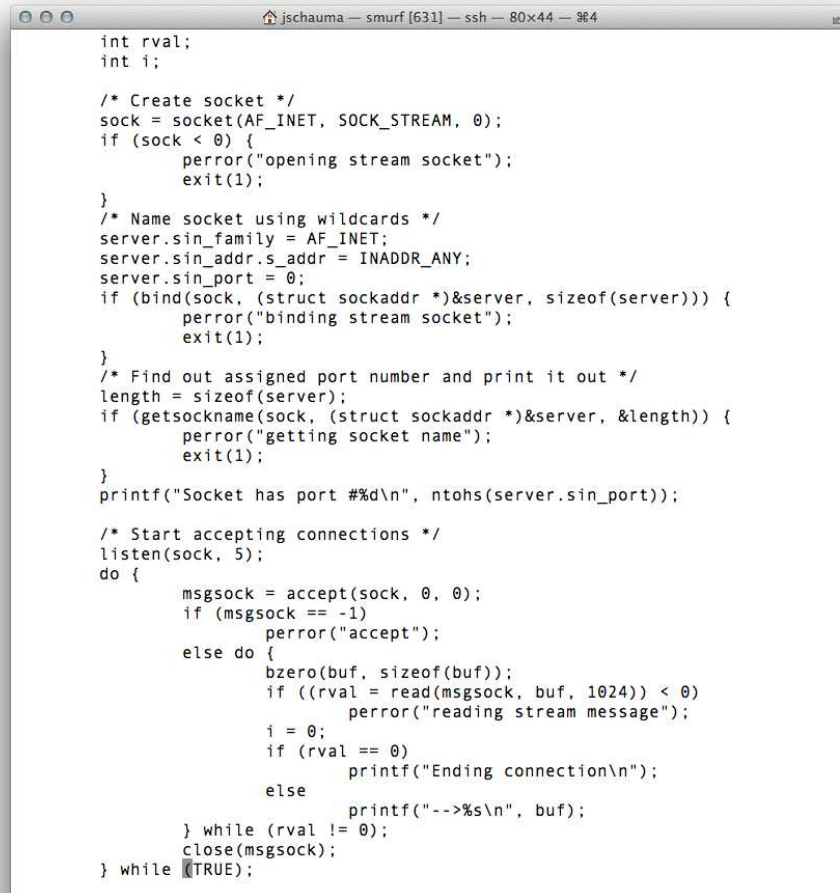
"Real World" edition

- group work:
 - 3-4 people in a team
 - let me know your team composition by next week
 - all people will get the same grade on the final project
- use of `git(1)` required
- use of a `Makefile` required
- code must be split across multiple files
- milestone deliverables (HW3 was the first)
- multi-platform support

Software Development Tools



Software Development Tools



```
int rval;
int i;

/* Create socket */
sock = socket(AF_INET, SOCK_STREAM, 0);
if (sock < 0) {
    perror("opening stream socket");
    exit(1);
}

/* Name socket using wildcards */
server.sin_family = AF_INET;
server.sin_addr.s_addr = INADDR_ANY;
server.sin_port = 0;
if (bind(sock, (struct sockaddr *)&server, sizeof(server))) {
    perror("binding stream socket");
    exit(1);
}

/* Find out assigned port number and print it out */
length = sizeof(server);
if (getsockname(sock, (struct sockaddr *)&server, &length)) {
    perror("getting socket name");
    exit(1);
}
printf("Socket has port %d\n", ntohs(server.sin_port));

/* Start accepting connections */
listen(sock, 5);
do {
    msgsock = accept(sock, 0, 0);
    if (msgsock == -1)
        perror("accept");
    else do {
        bzero(buf, sizeof(buf));
        if ((rval = read(msgsock, buf, 1024)) < 0)
            perror("reading stream message");
        i = 0;
        if (rval == 0)
            printf("Ending connection\n");
        else
            printf("-->%s\n", buf);
    } while (rval != 0);
    close(msgsock);
} while (TRUE);
```


Software Development Tools

UNIX Userland is an IDE – essential tools that follow the paradigm of “Do one thing, and do it right” can be combined.

The most important tools are:

- `$EDITOR`
- the compiler toolchain
- `gdb(1)` – debugging your code
- `make(1)` – project build management, maintain program dependencies
- `diff(1)` and `patch(1)` – report and apply differences between files
- `cvs(1)`, `svn(1)`, `git(1)` etc. – distributed project management, version control

Compilers



Compilers

A compiler translates *source code* from a high-level programming language into *machine code* for a given architecture by performing a number of steps:

- lexical analysis
- preprocessing
- parsing
- semantic analysis
- code generation
- code optimization

Compilers

There are many different closed- and open-source compiler chains:

- Intel C/C++ Compiler (or `icc`)
- Turbo C / Turbo C++ / C++Builder (Borland)
- Microsoft Visual C++
- ...

- Clang (a frontend to LLVM)
- GNU Compiler Collection (or `gcc`)
- Portable C Compiler (or `pcc`)
- ...

The compiler toolchain

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

Preprocessing

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

```
$ cd compilechain
$ cat hello.c
$ man cpp
$ cpp hello.c hello.i
$ file hello.i
$ man cc
$ cc -v -E hello.c > hello.i
$ more hello.i
$ cc -v -DFOOD=\"Avocado\" -E hello.c > hello.i.2
$ diff -bu hello.i hello.i.2
```

Compilation

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

```
$ more hello.i
$ cc -v -S hello.i > hello.s
$ file hello.s
$ more hello.s
```

Assembly

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

```
$ as -o hello.o hello.s
$ file hello.o
$ cc -v -c hello.s
$ objdump -d hello.o
[...]
```

Linking

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

```
$ ld hello.o
[...]  
$ ld hello.o -lc
[...]  
$ cc -v hello.o
[...]  
$ ld -dynamic-linker /lib64/ld-linux-x86-64.so.2 \  
    /usr/lib/x86_64-linux-gnu/crt1.o      \  
    /usr/lib/x86_64-linux-gnu/crti.o hello.o \  
    -lc /usr/lib/x86_64-linux-gnu/crtn.o  
$ file a.out  
$ ./a.out
```

Linking

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

```
$ cc -v -DFOOD=\"Avocado\" hello.c 2>&1 | more
```

`cc(1)` and `ld(1)`

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

Different flags can be passed to `cc(1)` to be passed through to each tool as well as to affect all tools.

```
$ cc -v -O2 -g hello.c 2>&1 | more
```

`cc(1)` and `ld(1)`

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

Different flags can be passed to `cc(1)` to be passed through to each tool as well as to affect all tools.

The order of the command line flags *may* play a role! Directories searched for libraries via `-L` and the resolving of undefined symbols via `-l` are examples of position sensitive flags.

```
$ cc -v main.c -L./lib2 -L./lib -lldtest 2>&1 | more
```

```
$ cc -v main.c -L./lib -L./lib2 -lldtest 2>&1 | more
```


`cc(1)` and `ld(1)`

The compiler usually performs preprocessing (via `cpp(1)`), compilation (`cc(1)`), assembly (`as(1)`) and linking (`ld(1)`).

Different flags can be passed to `cc(1)` to be passed through to each tool as well as to affect all tools.

The order of the command line flags *may* play a role! Directories searched for libraries via `-L` and the resolving of undefined symbols via `-l` are examples of position sensitive flags.

The behavior of the compiler toolchain may be influenced by environment variables (eg `TMPDIR`, `SGI_ABI`) and/or the compilers default configuration file (MIPSPro's `/etc/compiler.defaults` or gcc's `specs`).

```
$ cc -v hello.c
```

```
$ TMPDIR=/var/tmp cc -v hello.c
```

```
$ cc -dumpspec
```

A Debugger



`gdb(1)`

The purpose of a debugger such as `gdb(1)` is to allow you to see what is going on “inside” another program while it executes – or what another program was doing at the moment it crashed. `gdb` allows you to

- make your program stop on specified conditions (for example by setting *breakpoints*)
- examine what has happened, when your program has stopped (by looking at the *backtrace*, inspecting the value of certain variables)
- inspect control flow (for example by *stepping* through the program)

Other interesting things you can do:

- examine stack frames: *info frame*, *info locals*, *info args*
- examine memory: *x*
- examine assembly: *disassemble func*

gdb(1)

```
$ cd student
$ make
$ ./a.out -lR ~djd >/dev/null
total 2701553
[...]
Memory fault
$ gdb ./a.out
  run -lR ~djd
Starting program: /home/jschauma/t/student/a.out -lR ~djd
total 2701553
[...]
```

Program received signal SIGSEGV, Segmentation fault.

```
0x000000000040214a in print_entries (entryvect=0x6050a0, options=0x605010) at print.c:
221             printf("%10s ", grpentry->gr_name);
```

gdb(1)

```
(gdb) bt
#0  0x000000000040214a in print_entries (entryvect=0x6050a0, options=0x605010) at pri
#1  0x0000000000401ad5 in list_dir_contents (path=0x7fffffffec59 "/home/djd", ftsopts
#2  0x000000000040292e in main (argc=1, argv=0x7fffffffe9e8) at ls.c:57
(gdb) li
216     } else {
217         pwentry = getpwuid(myentry->fts_statp->st_uid);
218         printf("%8s ", pwentry->pw_name);
219
220         grpentry = getgrgid(myentry->fts_statp->st_gid);
221         printf("%10s ", grpentry->gr_name);
222     }
223
224     // displaying the size
225     size = (long long)(myentry->fts_statp->st_size);
(gdb) p grpentry
$1 = (struct group *) 0x0
```

`gdb(1)`

```
$ cc -g gdb1.c
$ ./a.out
[...]
$ gdb ./a.out
[...]
(gdb) break main
Breakpoint 1 at 0x400603: file gdb1.c, line 10.
(gdb) run
Starting program: /home/jschauma/t/gdb-examples/a.out

Breakpoint 1, main (argc=1, argv=0x7fffffffe9e8) at gdb1.c:10
10 c = fgetc(stdin);
(gdb) n
```

`gdb(1)`

```
$ ulimit -c unlimited
$ cc -g gdb2.c
$ ./a.out
$ gdb a.out core
bt
[...]
frame 2
li
p buf
kill
watch num
run
```

make(1)



make(1)

make(1) is a command generator and build utility. Using a description file (usually *Makefile*) it creates a sequence of commands for execution by the shell.

- used to sort out dependency relations among files

make(1)

`make(1)` is a command generator and build utility. Using a description file (usually *Makefile*) it creates a sequence of commands for execution by the shell.

- used to sort out dependency relations among files
- avoids having to rebuild the entire project after modification of a single source file

make(1)

`make(1)` is a command generator and build utility. Using a description file (usually *Makefile*) it creates a sequence of commands for execution by the shell.

- used to sort out dependency relations among files
- avoids having to rebuild the entire project after modification of a single source file
- performs *selective* rebuilds following a *dependency graph*

make(1)

`make(1)` is a command generator and build utility. Using a description file (usually *Makefile*) it creates a sequence of commands for execution by the shell.

- used to sort out dependency relations among files
- avoids having to rebuild the entire project after modification of a single source file
- performs *selective* rebuilds following a *dependency graph*
- allows simplification of rules through use of *macros* and *suffixes*, some of which are internally defined

make(1)

`make(1)` is a command generator and build utility. Using a description file (usually *Makefile*) it creates a sequence of commands for execution by the shell.


- used to sort out dependency relations among files
- avoids having to rebuild the entire project after modification of a single source file
- performs *selective* rebuilds following a *dependency graph*
- allows simplification of rules through use of *macros* and *suffixes*, some of which are internally defined
- different versions of `make(1)` (BSD make, GNU make, Sys V make, ...) may differ (among other things) in
 - variable assignment and expansion/substitution
 - including other files
 - flow control (for-loops, conditionals etc.)

make(1)

```
$ cd make-examples
```

```
$ ls *.*[ch]
```

cmp.c	ls.c	main.c	stat_flags.c	util.c
extern.h	ls.h	print.c	stat_flags.h	

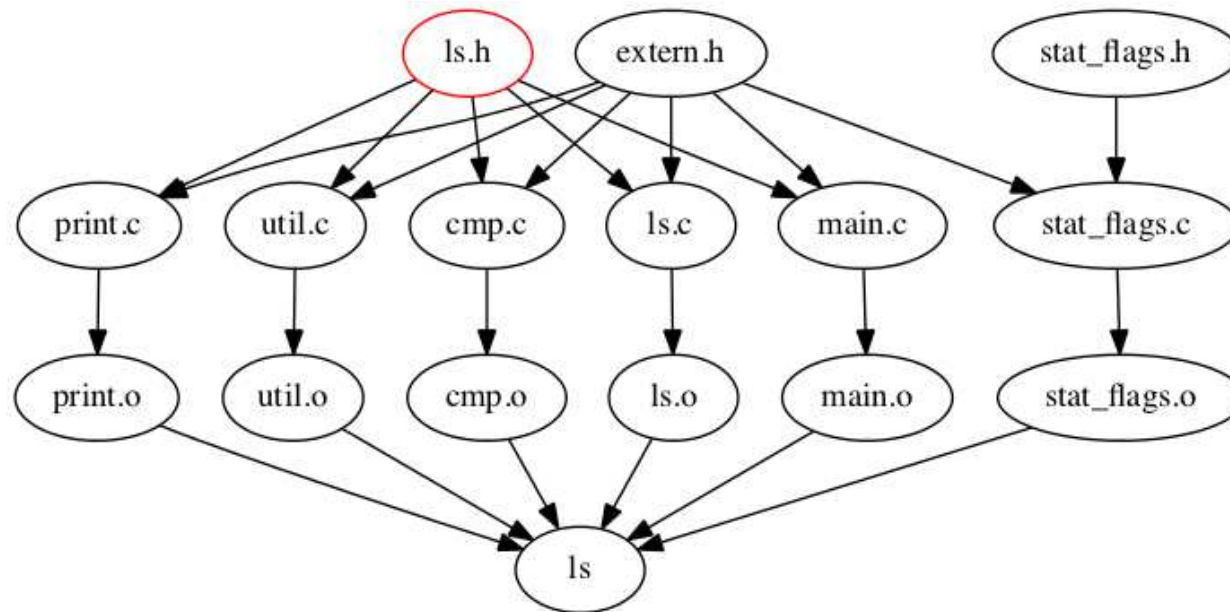


make(1)

```
$ cd make-examples
```

```
$ ls *.[ch]
```

cmp.c	ls.c	main.c	stat_flags.c	util.c
extern.h	ls.h	print.c	stat_flags.h	

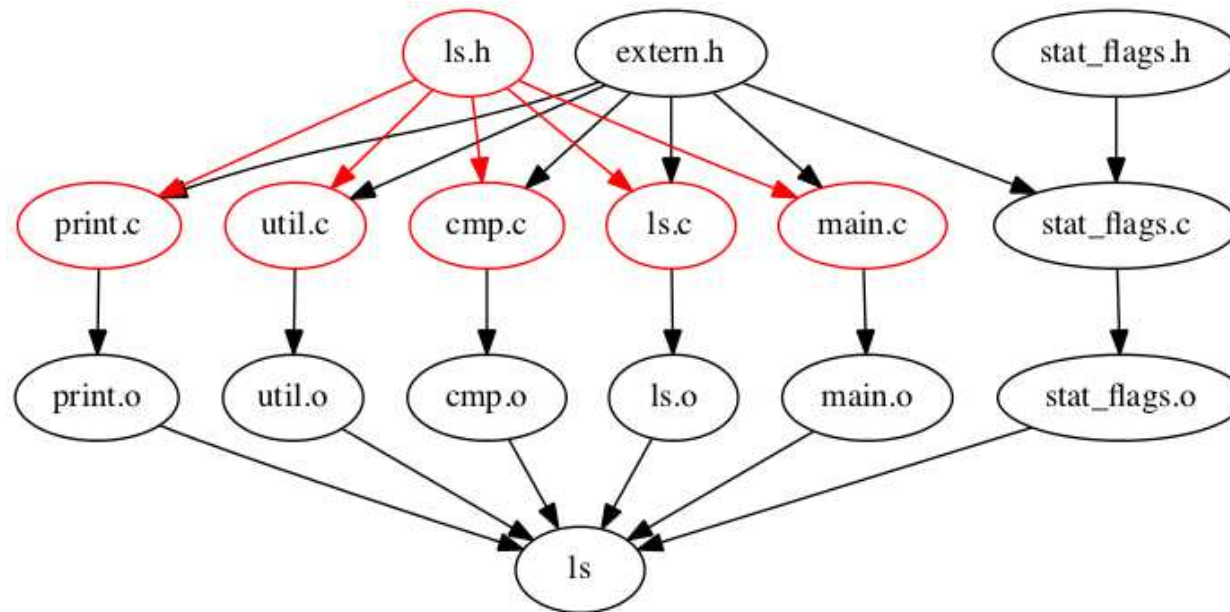


make(1)

```
$ cd make-examples
```

```
$ ls *. [ch]
```

cmp.c	ls.c	main.c	stat_flags.c	util.c
extern.h	ls.h	print.c	stat_flags.h	

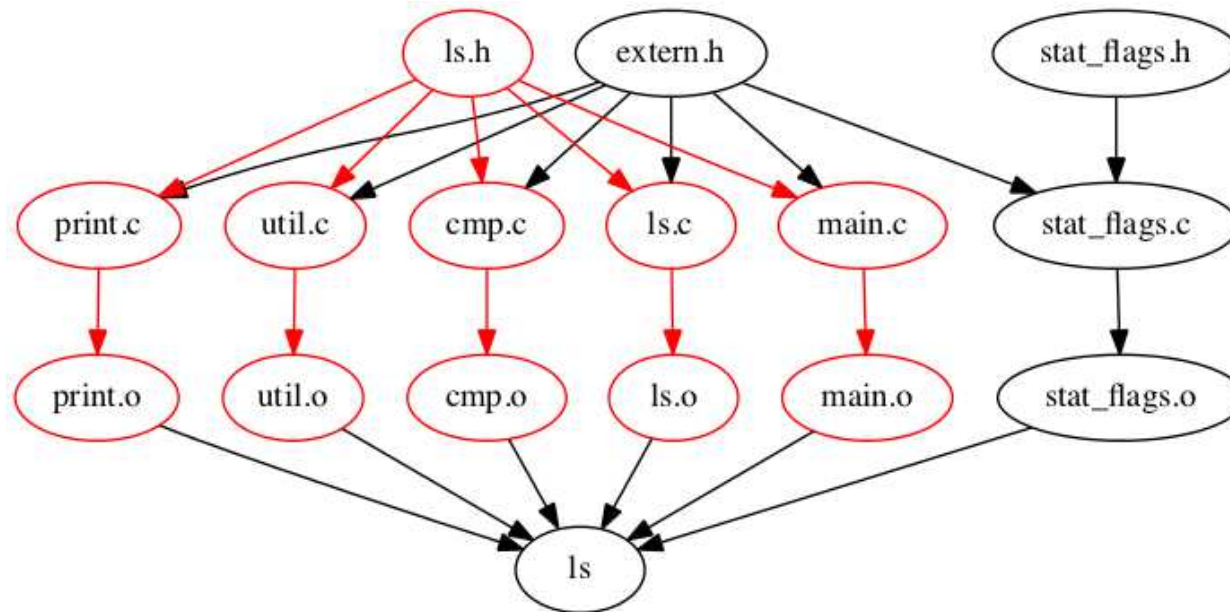


make(1)

```
$ cd make-examples
```

```
$ ls *. [ch]
```

cmp.c	ls.c	main.c	stat_flags.c	util.c
extern.h	ls.h	print.c	stat_flags.h	

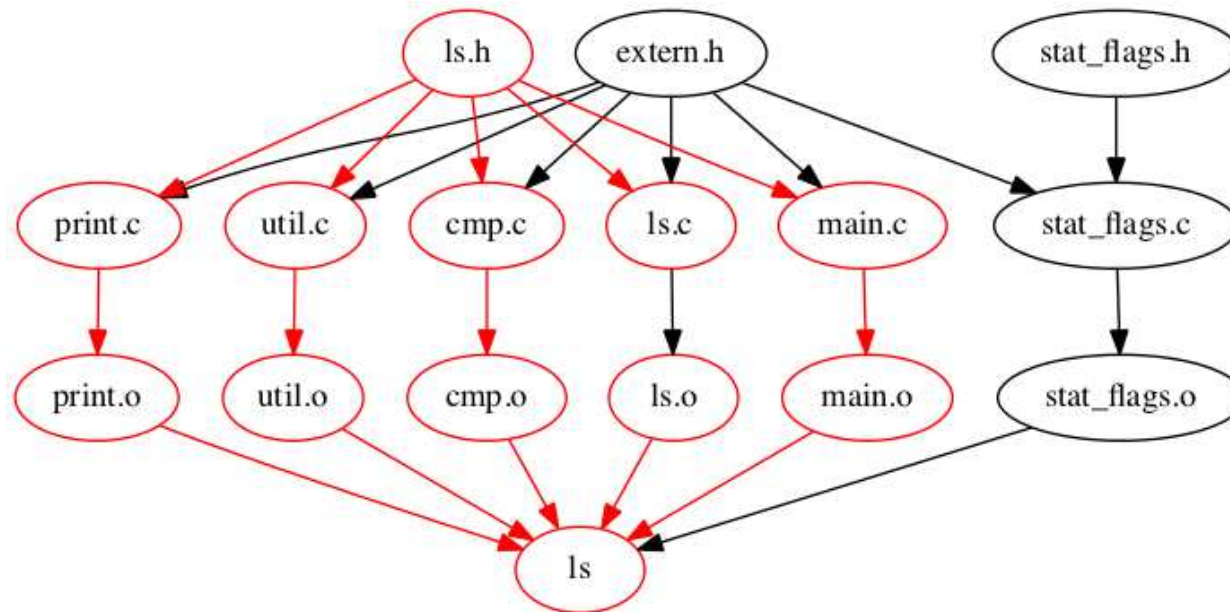


make(1)

```
$ cd make-examples
```

```
$ ls *.[ch]
```

cmp.c	ls.c	main.c	stat_flags.c	util.c
extern.h	ls.h	print.c	stat_flags.h	



make(1)

```
$ make -f Makefile.1
cc -c cmp.c
cc -c ls.c
cc -c main.c
cc -c print.c
cc -c stat_flags.c
cc -c util.c
cc cmp.o ls.o main.o print.o stat_flags.o util.o -o ls
$ touch ls.h
$ make -f Makefile.1
cc -c cmp.c
cc -c ls.c
cc -c main.c
cc -c print.c
cc -c util.c
cc cmp.o ls.o main.o print.o stat_flags.o util.o -o ls
$
```

make(1)

```
$ make -f Makefile.2
cc      -c -o cmp.o cmp.c
cc      -c -o ls.o ls.c
[...]
cc      -c -o util.o util.c
ls depends on cmp.o ls.o main.o print.o stat_flags.o util.o
cc cmp.o ls.o main.o print.o stat_flags.o util.o -o ls
$ bmake -f Makefile.2
gcc -pipe -O2 -c cmp.c
gcc -pipe -O2 -c ls.c
[...]
gcc -pipe -O2 -c util.c
ls depends on cmp.o ls.o main.o print.o stat_flags.o util.o
gcc cmp.o ls.o main.o print.o stat_flags.o util.o -o ls
$
```

make(1)

```
$ make -f Makefile.3
cc -c cmp.c -o cmp.o
cc -c ls.c -o ls.o
cc -c main.c -o main.o
cc -c print.c -o print.o
cc -Wall -g -c stat_flags.c -o stat_flags.o
cc -Wall -g -c util.c -o util.o
ls depends on cmp.o ls.o main.o print.o stat_flags.o util.o
cc cmp.o ls.o main.o print.o stat_flags.o util.o -o ls
$
```

Priority of Macro Assignments for `make(1)`

1. Internal (default) definitions of `make(1)`
2. Current shell environment variables. This includes macros that you enter on the *make* command line itself.
3. Macro definitions in *Makefile*.
4. Macros entered on the `make(1)` command line, if they follow the *make* command itself.

make(1)

```
$ bmake -f Makefile.4  
[...]  
$ bmake -f Makefile.2  
[...]  
$ CFLAGS=-Werror bmake -f Makefile.2  
[...]  
$ CFLAGS=-Werror bmake -f Makefile.4  
[...]  
$ bmake CFLAGS=-Werror -f Makefile.2  
[...]  
$ bmake CFLAGS=-Werror -f Makefile.4  
[...]
```

Ed is the standard text editor.

```
$ ed
?
help
?
quit
?
exit
?
bye
?
eat flaming death
?
^C
?
^D
?
```


Ed is the standard text editor.

```
$ ed
a
ed is the standard Unix text editor.
This is line number two.
.
2i

.
%1
3s/two/three/
w foo
q
$ cat foo
```

diff(1) and patch(1)

diff(1):

- compares files line by line
- output may be used to automatically edit a file
- can produce human “readable” output as well as diff entire directory structures
- output called a *patch*

diff(1) and patch(1)

patch(1):

- applies a `diff(1)` file (aka *patch*) to an original
- may back up original file
- may guess correct format
- ignores leading or trailing “garbage”
- allows for reversing the patch
- may even correct context line numbers

diff(1) and patch(1)

```
$ diff Makefile.2 Makefile.4
8c8
< #CFLAGS= -Wall -g
---
> CFLAGS= -Wall -g
$ cp Makefile.2 /tmp
$ ( diff -e Makefile.2 Makefile.4; echo w; ) | ed Makefile.2
$ diff Makefile.[24]
$ mv /tmp/Makefile.2 .
$ diff -c Makefile.[24]
$ diff -u Makefile.[24] > /tmp/patch
$ patch </tmp/patch
$ diff Makefile.[24]
```

Revision Control

Version control systems allow you to

- collaborate with others
- simultaneously work on a code base
- keep old versions of files
- keep a log of the who, when, what, and why of any changes
- perform release engineering by creating *branches*

Revision Control

- Source Code Control System (*SSCS*) begat the Revision Control System (*RCS*).
- RCS operates on a single file; still in use for misc. OS config files
- the Concurrent Versions System (*CVS*) introduces a client-server architecture, control of hierarchies
- *Subversion* provides atomic commits, renaming, cheap branching etc.
- *Git*, *Mercurial* etc. implement a *distributed* approach (ie peer-to-peer versus client-server), adding other features (cryptographic authentication of history, ...)

Revision Control

Examples:

`http://cvsweb.netbsd.org/bsdweb.cgi/src/bin/ls/`

`http://svnweb.freebsd.org/base/stable/9/bin/ls/`

`http://git.savannah.gnu.org/cgit/coreutils.git/log/`

Setting up git(1) for your final project

Each team identifies one member to host the repository.

Initialize the source code repository:

```
ssh linux-lab.cs.stevens.edu
mkdir -p ~/git/sws.git
cd ~/git/sws
git --bare init
setfacl -R -m user:teammate1:rwx .
setfacl -R -m user:teammate2:rwx .
setfacl -R -m user:teammate3:rwx .
```


Setting up git(1) for your final project

Add your code from the project directory:

```
ssh linux-lab.cs.stevens.edu
cd ~/cs631/sws
git init
git remote add origin ~/git/sws.git
git config branch.master.merge refs/heads/master
git config branch.master.remote origin
git config --global user.name "Your Full Name"
git config --global user.email $USER@stevens.edu
git add *
git commit -m "initial import of all files"
git push origin master
```

Setting up git(1) for your final project

Check out a copy on a local host and make changes:

```
local$ git clone linux-lab.cs.stevens.edu:git/sws.git
local$ cd sws
local$ $EDITOR README
local$ git add README
local$ git commit README
local$ git push
```

```
linux-lab$ cd ~/cs631/sws
linux-lab$ git pull
linux-lab$ $EDITOR ...
linux-lab$ git add ...; git commit; git push
```

Homework

Generate an ssh key pair and send me the *public* key.
See `ssh-keygen(1)`.

Links

Revision Control:

<http://cvsbook.red-bean.com/cvsbook.html>

<http://svnbook.red-bean.com/>

<http://git-scm.com/>

GDB:

http://sources.redhat.com/gdb/current/onlinedocs/gdb_toc.html

<http://heather.cs.ucdavis.edu/~matloff/UnixAndC/CLanguage/Debug.html>

<http://www.unknownroad.com/rtfm/gdbtut/gdbtoc.html>